Snapshot of October 29, 2026. This dossier describes the four repositories as they stood that day; the projects have moved on since, and the repositories are the authority. Underlined links lead to the sources, frozen at their state of that day.

VERIFY

Verify a figure yourself

Take the most visible figure on the site: “2.12× faster”. In five steps, without installing anything, you will find it in the repository’s raw receipt — and see where it comes from.

1. The figure

At the top of the gnostral.rs dossier, the first box says: **“2.12× faster than the previous version of the same engine”**. That is a claim. The question now is what it rests on.

2. The milestone

In the “What was measured” table, the *Borrowed experts* row carries the label **Q-010**. Hovering over it opens a card that leads to the measurement’s audit receipt. Every figure in the dossier has a milestone, and every milestone a receipt.

3. The receipt, at the frozen commit

The link opens this file, in the public repository:

```
github.com/gnu6dotlive/gnostral.rs/blob/38c4134a3a398b5b83d34326826362cf4dcc7336/
  evidence/runs/q010-hotpath-ssd-ab-audit-20261010.json
```

The long string of digits and letters is the **commit** fingerprint: it names one exact version of the repository. The file read at this address cannot have changed since October 29, even though the repository has moved on.

4. The three fields

The receipt is long, but three fields are enough. They are keyed by answer length; the site's figure is for 512-token answers:

```
"aggregate": {
  "baseline":  { "median_tps": { "512": 9.506482 } },
  "candidate": { "median_tps": { "512": 20.179588 } }
},
"candidate_over_baseline_median_tps_ratios": { "512": 2.122718793345425 }
```

baseline is the engine without the patch, *candidate* the same engine with it. *median_tps* is the median throughput, in tokens per second, over three processes.

5. Redo the arithmetic

20.179588 ÷ 9.506482 = **2.1227…**, rounded to **2.12**. That is exactly the ratio the receipt itself writes. You can go one step further: the *per_request* field lists the twelve answers one by one, and the medians can be found by hand (9.51 among 10.12 · 9.51 · 9.51; 20.18 among 20.18 · 20.21 · 19.48).

> THE SAME CHECK, IN ONE TERMINAL LINE

What you just proved — and what you did not

Proved

The site's figure is the receipt's, at a dated commit. The receipt is consistent with itself: the ratio follows from the medians, and the medians from the twelve answers.

Not proved

That the measurement happened as described. That would take the machine and the protocol. The receipt gives what you need to rebuild them — fingerprints of the model and binary, thresholds declared in advance, the failed first run — but that is another job.